

Generate Better AI Images in Minutes

The Prompt Engineering Toolkit

A practical guide for agencies and creative teams

6 AI Models · Optimized Prompts · Zero Guesswork

inspiredcreativegroupinc.com

What's Inside This Toolkit

This guide and the **Image Prompt Engineer** skill work together as a complete toolkit.

■ The Skill (included in your download)

A Claude skill that transforms plain English descriptions into production-ready, model-optimized prompts for 6 leading AI image generators. It handles model selection, keyword optimization, quality scoring, and technical specs — so you don't have to.

■ This Guide

Tells you *which* model to use, *when*, and *why*. Covers the strengths of each model, where to generate images (from free web UIs to API pipelines), and 5 prompt engineering techniques that make an immediate difference in output quality.

How they work together:

1. You describe the image you want in plain English
2. The skill optimizes your prompt for the best model
3. You take the output to any generation service — web UI or API
4. This guide helps you understand which model and service fits your workflow

The skill does **not** generate images itself — it engineers the best possible prompt. This means you're never locked into a single provider. Use whichever service works for you.

Why These 6 Models

We evaluated dozens of AI image generators across photorealism, text rendering, design quality, speed, resolution, editing capabilities, and price. These 6 cover every professional use case with no gaps and no redundancy:

Model	Covers
Flux 2	Photorealism, multi-reference consistency, detailed scenes
Nano Banana Pro	High-resolution (up to 4K), general purpose, natural language
Recraft V4	Vector/SVG output, design assets, marketing materials
Ideogram 3.0	Text rendering, color-controlled branding, style codes
GPT Image 1.5	Instruction following, conversational editing, iteration
Grok Imagine	Quick aesthetics, budget batch work, phone-screen ratios

Selection criteria: Each model is best-in-class at something no other model does as well. Flux 2 owns photorealism. Recraft V4 is the only model producing real SVG vectors. Ideogram 3.0 leads in text rendering. GPT Image 1.5 is uniquely conversational. Together, they give you full coverage — from quick \$0.02 drafts to production 4K prints.

The 6 Models

MODEL 1 OF 6

Flux 2

by Black Forest Labs

Best for: Photorealistic product shots, portraits, detailed scenes, multi-reference consistency

Strengths

- Best photorealism of any current model — natural skin, lighting, materials
- 32K token prompt context — more detail = better results
- Up to 10 reference images for character/style consistency across a campaign
- Sub-10-second generation (Pro/Flex variants)
- Strong text rendering, especially with the Flex variant

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- Leonardo.ai — paste prompt, select Flux model
- Freepik AI Image Generator
- Any interface offering Flux 2 models

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- Fal.ai — fal-ai/flux-2-flex (~\$0.025/image)
- BFL.ai — direct from Black Forest Labs (~\$0.025/image)
- Replicate — flux-2-flex (usage-based pricing)

■ **Price range:** \$0.025–\$0.05/image (varies by variant and provider)

■ **Pro tip:** *Enable prompt expansion (on by default) — the model enhances your prompt with its own knowledge. For product photography, provide reference images for consistency across an entire catalog.*

MODEL 2 OF 6

Nano Banana Pro (Gemini 3 Pro Image)

by Google

Best for: High-resolution output (up to 4K), general purpose, natural language prompting

Strengths

- Up to 4K resolution — ideal for large-format prints and gallery work
- Simple, conversational prompting — describe naturally, no keyword engineering needed
- Google's multimodal intelligence understands context and composition
- Great all-rounder for teams that want one reliable default model

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- Google AI Studio (Gemini) — native access
- Any Gemini-powered creative tool

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- Fal.ai — fal-ai/nano-banana-pro (~\$0.02/image)
- Google AI API — direct Gemini access

■ Price range: ~\$0.02/image

■ Pro tip: Write prompts like you're describing the image to a friend — conversational language outperforms keyword stacking with this model. Use '4K' resolution tier for anything going to print.

MODEL 3 OF 6

Recraft V4

by Recraft

Best for: Logos, icons, vector graphics (SVG), design assets, marketing materials

Strengths

- The ONLY model producing real, editable SVG vector files
- Outputs: SVG, PNG, JPG, PDF, TIFF, and even Lottie animations
- Design-forward aesthetic — understands visual hierarchy and composition
- Exploration mode gives multiple design directions from one prompt

- Free plan available on [recraft.ai](#)

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- Recraft.ai — free web interface with full vector support
- Any design tool with Recraft integration

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- Fal.ai — [fal-ai/reCraft-v4](#) (pricing TBD)
- Recraft.ai API — direct access (free plan available)
- Replicate — [reCraft-v4](#) (usage-based pricing)

■ **Price range:** Free plan available; API pricing varies

■ **Pro tip:** Use design vocabulary in your prompts — 'balanced composition,' 'cohesive color palette,' 'visual hierarchy.' Recraft was built for designers and responds to design language better than photography terms.

MODEL 4 OF 6

Ideogram 3.0

by Ideogram

Best for: Text-in-image rendering, branded content, color-controlled designs, reproducible styles

Strengths

- Industry gold standard for text rendering — signs, logos, titles, dense typography
- Style codes (8-char hex) for exact style reproduction across campaigns
- Color palette with hex weights for brand-precise color control
- Negative prompts to explicitly exclude unwanted elements
- Magic prompt auto-enhancement for quick generation

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- Ideogram.ai — free web interface, best text rendering UX
- ChatGPT (with DALL-E) — for text-heavy requests
- Any creative tool with Ideogram integration

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- Fal.ai — fal-ai/ideogram-3 (~\$0.02/image)
- Ideogram.ai API — direct access
- Together.ai — ideogram/ideogram-3.0
- Replicate — ideogram-3.0

■ **Price range:** ~\$0.02/image; free tier available on ideogram.ai

■ **Pro tip:** Always put text you want rendered in quotes within the prompt — 'SALE 50% OFF'. Save your style codes and reuse them across an entire campaign for perfect visual consistency.

MODEL 5 OF 6

GPT Image 1.5

by OpenAI

Best for: Instruction following, multi-turn editing, conversational refinement, concept iteration

Strengths

- Best instruction following — understands complex, nuanced directions
- Unique multi-turn editing: generate, then refine ('make the sky more orange')
- Broad real-world knowledge — can reference brands, locations, styles accurately
- Three quality tiers: low (\$0.009), medium, high (\$0.20) — budget flexibility
- Transparent background support

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- ChatGPT — the native interface, supports multi-turn editing
- OpenAI Playground — direct access with parameter control

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- OpenAI API — gpt-image-1.5 (\$0.009–\$0.20/image depending on quality)

■ Price range: \$0.009 (low) to \$0.20 (high) per image

■ Pro tip: *Structure your first prompt with clearly separate elements — this makes follow-up edits much easier. 'A woman in a red dress, standing in a modern office, with city skyline through windows' lets you later say 'change the dress to blue' without regenerating everything.*

MODEL 6 OF 6

Grok Imagine

by xAI

Best for: Quick aesthetic images, budget-friendly batch work, phone wallpapers, exploration

Strengths

- Cheapest option at \$0.02/image — great for exploration and drafts
- Unique phone-screen aspect ratios (19.5:9, 20:9) for mobile content
- Fast generation with attractive aesthetic output
- Combined image generation + editing in one API family
- Simple API with minimal configuration needed

■■ Where to Generate (Copy-Paste)

Paste the optimized prompt directly into any of these web interfaces:

- Grok (x.ai / X platform) — native interface
- Any tool with Grok Imagine integration

■ API Providers (For Automation)

Connect via API for batch processing, pipelines, and agent workflows:

- Fal.ai — xai/grok-imagine-image (\$0.02/image)
- xAI API — direct access

■ **Price range:** ~\$0.02/image

■ **Pro tip:** *Focus on mood and aesthetics rather than technical details — this model interprets artistic intent better than precise specifications. Perfect for quick mood boards before committing to a more detailed generation with Flux 2 or Ideogram.*

Comparison Matrix

	Flux 2	Nano Banana	Recraft V4	Ideogram 3.0	GPT Image	Grok Imagine
Photorealism	★★★★★	★★★★	★★★★	★★★★	★★★★	★★★
Text Render	★★★★★	★★★	★★★★	★★★★★	★★★★★	★★★
Speed	★★★★★	★★★★	★★★	★★★★	★★★	★★★★★
Max Res.	4MP	4K	2048 ²	Varies	1536px	Varies
Vector/SVG	—	—	✓	—	—	—
Editing	★★★★	★	★	★★★★	★★★★★	★★★
Price/img	\$0.025+	~\$0.02	Free+	~\$0.02	\$0.009+	~\$0.02

Quick Decision Guide

Need a realistic photo? → Flux 2

Need text in the image? → Ideogram 3.0 or GPT Image 1.5

Need a logo or vector? → Recraft V4

Need 4K for print? → Nano Banana Pro

Need to iterate/edit? → GPT Image 1.5

Need it fast and cheap? → Grok Imagine (\$0.02)

Need brand consistency? → Flux 2 (references) or Ideogram (style codes)

Need a phone wallpaper? → Grok Imagine (19.5:9 ratio)

5 Prompt Tips That Actually Work

1. Be Specific About What Makes It Good

Don't say 'beautiful sunset.' Say 'golden hour sunset with warm amber light casting long shadows across wet sand, scattered clouds lit from below in pink and coral, gentle wave foam in the foreground.' The more specific your description, the less the AI has to guess — and guessing is where quality drops.

2. Front-Load the Subject

Put your main subject first, then add context. 'A golden retriever sitting in a sunlit meadow, wildflowers in soft focus behind, gentle breeze, photorealistic' works better than burying the subject mid-sentence. AI models weight the beginning of prompts more heavily.

3. Specify Lighting — Always

Lighting is the single biggest quality lever in AI image generation. 'Studio lighting with soft shadows' produces dramatically better results than leaving it unspecified. Good options: golden hour, blue hour, studio lighting, natural window light, backlit, dramatic rim lighting, overcast soft light.

4. Match Your Language to the Model

Each model responds to different prompt styles. Flux 2 rewards detail and photography terms. Nano Banana Pro and GPT Image 1.5 prefer natural conversation. Recraft V4 responds to design vocabulary. Grok Imagine wants aesthetic and mood keywords. The skill handles this automatically, but knowing it helps you write better starting descriptions.

5. Use the Negative (When Available)

Ideogram 3.0 supports negative prompts — telling the AI what to avoid. 'Negative: blurry, low quality, distorted faces, cluttered background' can prevent common failure modes. Even for models without formal negative prompts, you can add 'avoid blurry details, no clutter' to the end of your positive prompt.

How to Use the Skill

The Image Prompt Engineer skill works in any Claude environment. Here's how to get started:

Claude Code (Terminal)

Drop the skill folder into your workspace's skills directory. Claude automatically picks it up — just start describing images.

Claude Desktop

Add the skill folder as a project in Claude Desktop settings. Include SKILL.md and the references/ folder. Start a conversation and describe what you want.

Claude Web (claude.ai)

Upload SKILL.md as an attachment in a new conversation. For deeper model knowledge, also upload the files from the references/ folder. Tell Claude: 'Use the attached skill to optimize my image prompts.'

What You'll Get

For every image description, the skill returns:

- An **optimized prompt** — human-readable, copy-paste ready
- A **model recommendation** — which AI model fits your needs
- **Provider options** — where to generate, with pricing
- A **quality score** (0–10) — how effective the prompt is
- **Technical specs** — resolution, aspect ratio, parameters
- **Warnings** — anything that might cause issues

The optimized prompt text is always written in plain English — you can paste it directly into any image generation interface, from Midjourney to Canva AI to any API.

Ready to See What AI-Powered Production Looks Like at Scale?

This toolkit is one piece of a 40+ agent system we built for creative production — video, music, design, and strategy, all working together.

We don't compete with your agency — we complete it. Whether you need production now or want to build AI workflows for your own team, let's talk.

[Schedule a Discovery Call](#)

inspiredcreativegroupinc.com/partners