
One Skill From Our Pipeline. Yours Free.

The Prompt Engineering Toolkit

A practical guide for agencies and creative teams.

6 AI Models · Film Stocks & Lenses · Anti-AI Realism

inspiredcreativegroupinc.com

WHAT'S INSIDE

This Toolkit

This guide and the **Image Prompt Engineer** skill work together. The skill teaches AI to think like a photographer — the guide tells you what to aim at.

■ The Skill	■ This Guide
Transforms plain English into production-ready prompts using real photography knowledge — film stocks, lens focal lengths, lighting setups, and anti-AI realism techniques mapped across 6 models. Includes quality scoring and per-model optimization.	Covers which model to use, when, and why. Where to generate — from free web UIs to API pipelines. The photography intelligence that powers the skill. Plus techniques that make an immediate difference.

How they work together:

1. Describe the image you want in plain English.
2. The skill optimizes your prompt for the best model.
3. Take the output to any generation service — web UI or API.
4. This guide helps you pick the right model and service for your workflow.

The skill does **not** generate images — it engineers the best possible prompt. You're never locked into a single provider.

MODEL SELECTION

Why These 6 Models

We evaluated dozens of AI image generators across photorealism, text rendering, design quality, speed, resolution, editing, and price. These 6 cover every professional use case — no gaps, no redundancy.

Model	What It Covers
Flux 2	Photorealism, multi-reference consistency, detailed scenes
Nano Banana Pro	High-resolution (up to 4K), general purpose, natural language
Recraft V4	Vector/SVG output, design assets, marketing materials
Ideogram 3.0	Text rendering, color-controlled branding, style codes
GPT Image 1.5	Instruction following, conversational editing, iteration
Grok Imagine	Quick aesthetics, budget batch work, phone-screen ratios

Each model is best-in-class at something no other does as well. Flux 2 owns photorealism. Recraft V4 is the only model producing real SVG vectors. Ideogram 3.0 leads in text rendering. GPT Image 1.5 is uniquely conversational. Together: full coverage from \$0.02 drafts to production 4K prints.



THE HEADLINE FEATURE

Camera & Photography Intelligence

The same photography knowledge professionals spend years learning, mapped to 6 AI image models through empirical testing.

100mm macro · Kodak Portra 400 · Rembrandt lighting

The Photography Knowledge Base

This isn't prompt templates — it's real photography knowledge mapped to 6 AI image models through empirical testing.

Film Stock Color Science

Each film stock shifts AI color rendering in specific, repeatable ways. The skill maps these to model-specific behavior:

Film Stock	Character
Kodak Portra 400	Warm skin tones, creamy highlights, subtle grain
CineStill 800T	Tungsten-balanced, halation glow, neon color shift
Fuji Pro 400H	Cool greens, muted pastels, soft contrast
Kodachrome 64	Punchy saturation, deep reds, dense shadows
Ilford HP5	Classic B&W, rich midtones, visible grain structure

Lens Focal Length → Composition

- 24mm wide-angle → environmental context, dramatic perspective
- 50mm standard → natural perspective, documentary feel
- 85mm portrait → flattering compression, creamy background separation
- 135mm telephoto → strong compression, isolated subjects, cinematic depth

Camera Body Aesthetic Anchors

Naming a camera body shifts the AI's entire rendering approach:

- “Hasselblad X2D” → medium format creaminess, extreme detail
- “Leica M6” → street photography character, rangefinder aesthetic
- “Fujifilm X-T5” → warm film-like tones, rich color science

Professional Lighting Library

Lighting is the most reliable parameter across ALL 6 models. The skill includes a full library:

- Rembrandt — triangle highlight on cheek, dramatic portraiture
- Butterfly — glamour/beauty, overhead with fill below
- Split — half-face illumination, maximum drama
- Rim / edge — silhouette separation, cinematic depth
- Golden hour — the universal quality booster across every model
- Practical — in-scene light sources (lamps, neon, candles) for authenticity

Every parameter includes a **per-model compatibility matrix** (■/■/■) showing exactly what works on which model — so you never waste a generation on an unsupported feature.

Anti-AI Realism System

The biggest myth in AI image generation: “8K ultra-detailed hyper-realistic” makes images look better. The truth? **It makes them look more artificial.**

Real photographs have imperfections. AI images are too perfect — that’s how you spot them. The skill includes per-model realism recipes based on empirical testing:

The 4 Realism Levers

- **Imperfection cues** — film grain, uneven lighting, slight chromatic aberration, dust particles. The #1 realism lever across all models.
- **Capture context** — “shot on location” beats “professional studio.” Describing how/where the photo was taken adds authenticity that equipment names don’t.
- **Texture-first approach** — prioritize material textures (weathered wood, matte skin, rough concrete) over global quality keywords. Specific textures = specific realism.
- **What-to-avoid guidance** — per-model anti-patterns that waste credits. No negative prompts on Flux 2. Don’t ignore Ideogram’s style mode selector. GPT Image’s “studio polish” trap. Each model has specific pitfalls.

The skill includes **copy-paste realism recipes** for each of the 6 models — tested combinations that consistently produce images indistinguishable from real photographs.

■ *Based on empirical testing — Reddit A/B comparisons, official documentation, and community validation. Not guesswork.*

6 Models, Explained

MODEL 1 OF 6

Flux 2

by Black Forest Labs

Best for: Photorealistic product shots, portraits, detailed scenes, multi-reference consistency

Key Strengths

- Best photorealism of any current model — natural skin, lighting, materials
- 32K token prompt context — more detail always helps
- Up to 10 reference images for character/style consistency across campaigns
- Sub-10-second generation (Pro/Flex variants)
- Strong text rendering, especially with the Flex variant

Copy-Paste (Web UI)	API (Automation)
Leonardo.ai	Fal.ai — ~\$0.025/image
Freepik AI Image Generator	BFL.ai (direct) — ~\$0.025/image
Any interface offering Flux 2	Replicate — usage-based

Price: \$0.025–\$0.05/image depending on variant and provider

■ *Enable prompt expansion (on by default) — the model enhances your prompt with its own knowledge. For product photography, provide reference images for consistency across an entire catalog.*

MODEL 2 OF 6

Nano Banana Pro

by Google (Gemini 3 Pro Image)

Best for: High-resolution output (up to 4K), general purpose, natural language prompting

Key Strengths

- Up to 4K resolution — ideal for large-format prints and gallery work
- Simple, conversational prompting — describe naturally, skip the keywords
- Google's multimodal intelligence understands context and composition
- Great all-rounder for teams that want one reliable default model

Copy-Paste (Web UI)	API (Automation)
Google AI Studio (Gemini)	Fal.ai — ~\$0.02/image
Any Gemini-powered creative tool	Google AI API — direct Gemini access

Price: ~\$0.02/image

■ *Write prompts like you're describing the image to a colleague — conversational language outperforms keyword stacking with this model. Use the '4K' resolution tier for anything going to print.*

MODEL 3 OF 6

Recraft V4

by Recraft

Best for: Logos, icons, vector graphics (SVG), design assets, marketing materials

Key Strengths

- The ONLY model producing real, editable SVG vector files
- Outputs: SVG, PNG, JPG, PDF, TIFF, and Lottie animations
- Design-forward aesthetic — understands hierarchy and composition
- Exploration mode: multiple design directions from one prompt
- Free plan available on recraft.ai

Copy-Paste (Web UI)	API (Automation)
Recraft.ai — free web interface	Fal.ai — pricing TBD
Any tool with Recraft integration	Recraft.ai API — free plan available
	Replicate — usage-based

Price: Free plan available; API pricing varies

■ *Use design vocabulary: ‘balanced composition,’ ‘cohesive color palette,’ ‘visual hierarchy.’ Recraft was built for designers and responds to design language better than photography terms.*

MODEL 4 OF 6

Ideogram 3.0

by Ideogram

Best for: Text-in-image rendering, branded content, color-controlled designs, reproducible styles

Key Strengths

- Gold standard for text rendering — signs, logos, titles, dense typography
- Style codes (8-char hex) for exact style reproduction across campaigns
- Color palette with hex weights for brand-precise color control
- Negative prompts to explicitly exclude unwanted elements
- Magic prompt auto-enhancement for quick generation

Copy-Paste (Web UI)	API (Automation)
Ideogram.ai — free web interface	Fal.ai — ~\$0.02/image
ChatGPT (for text-heavy requests)	Ideogram.ai API — direct
Any tool with Ideogram integration	Together.ai
	Replicate

Price: ~\$0.02/image; free tier on ideogram.ai

■ Always put text you want rendered in quotes within the prompt — “SALE 50% OFF”. Save your style codes and reuse them across an entire campaign for perfect visual consistency.

MODEL 5 OF 6

GPT Image 1.5

by OpenAI

Best for: Instruction following, multi-turn editing, conversational refinement, concept iteration

Key Strengths

- Best instruction following — understands complex, nuanced directions
- Unique multi-turn editing: generate, then refine ('make the sky more orange')
- Broad real-world knowledge — references brands, locations, styles accurately
- Three quality tiers: low (\$0.009), medium, high (\$0.20) — budget flexibility
- Transparent background support

Copy-Paste (Web UI)	API (Automation)
ChatGPT — native, multi-turn editing	OpenAI API — \$0.009–\$0.20/image
OpenAI Playground	

Price: \$0.009 (low) to \$0.20 (high) per image

■ *Structure your first prompt with clearly separate elements — this makes follow-up edits easier. 'A woman in a red dress, standing in a modern office, with city skyline through windows' lets you later say 'change the dress to blue' without regenerating everything.*

MODEL 6 OF 6

Grok Imagine

by xAI

Best for: Quick aesthetic images, budget batch work, phone wallpapers, exploration

Key Strengths

- Cheapest option at \$0.02/image — great for exploration and drafts
- Unique phone-screen ratios (19.5:9, 20:9) for mobile content
- Fast generation with attractive aesthetic output
- Combined image generation + editing in one API family

Copy-Paste (Web UI)	API (Automation)
Grok (x.ai / X platform)	Fal.ai — \$0.02/image
Any tool with Grok integration	xAI API — direct access

Price: ~\$0.02/image

■ Focus on mood and aesthetics rather than technical details. Perfect for quick mood boards before committing to a detailed generation with Flux 2 or Ideogram.

AT A GLANCE

Comparison Matrix

	Flux 2	Nano Banana	Recraft V4	Ideogram 3.0	GPT Image	Grok Imagine
Photo-realism	★★★★★	★★★★	★★★★	★★★★	★★★★	★★★
Text	★★★★★	★★★	★★★★	★★★★★	★★★★★	★★★
Speed	★★★★★	★★★★	★★★	★★★★	★★★	★★★★★
Max Res	4MP	4K	2048 ²	Varies	1536px	Varies
SVG	—	—	✓	—	—	—
Editing	★★★★	★	★	★★★★	★★★★★	★★★
Price	\$0.025+	~\$0.02	Free+	~\$0.02	\$0.009+	~\$0.02

Quick Decision Guide

Need a realistic photo? → Flux 2

Need text in the image? → Ideogram 3.0 or GPT Image 1.5

Need a logo or vector? → Recraft V4

Need 4K for print? → Nano Banana Pro

Need to iterate and edit? → GPT Image 1.5

Need it fast and cheap? → Grok Imagine (\$0.02)

Need brand consistency? → Flux 2 (references) or Ideogram (style codes)

TECHNIQUES

5 Techniques That Actually Work

Fuji Pro 400H · 50mm Summicron · natural window light



5 Techniques That Actually Work

1 Think Like a Photographer, Not a Prompter

Instead of ‘beautiful photo of a woman,’ think about what a photographer would set up: ‘Rembrandt lighting, 85mm f/1.4, Kodak Portra 400 film stock, shallow depth of field.’ Real photography parameters produce real-looking results. The skill maps these automatically.

2 Specify Lighting — The #1 Lever

Lighting is the single most reliable parameter across ALL 6 models. ‘Golden hour side-lighting with long shadows’ produces dramatically better results than leaving it unspecified. The skill includes a full lighting library — Rembrandt, butterfly, split, rim, practical — each mapped per model.

3 Add Imperfections for Realism

The fastest way to make AI images look less AI: add film grain, slight chromatic aberration, dust motes, uneven natural lighting. Real photos aren’t perfect. ‘Shot on Fuji Pro 400H, slight vignette, natural window light’ beats ‘ultra-realistic 8K.’

4 Stop Using ‘8K Ultra-Detailed’

Quality keyword spam (‘8K, ultra-detailed, hyper-realistic, masterpiece’) is the most common mistake. It actually makes images look MORE artificial by pushing toward synthetic perfection. The skill replaces these with empirically-tested parameters that work per model.

5 Match Language to the Model

Each model responds differently. Flux 2 rewards photography terms and film stocks. Nano Banana Pro and GPT Image prefer conversational descriptions. Recraft V4 responds to design vocabulary. Grok’s ‘Aggressive Realism’ technique is model-specific. The skill handles all of this automatically.

GETTING STARTED

How to Use the Skill

This skill is a plain markdown file. It works with any AI assistant that can read uploaded documents — not just Claude. Here's how to set it up in each environment:

Claude Code (Terminal)

Place the skill folder in your workspace's skills/ directory. Claude discovers it automatically and uses it whenever you describe an image.

→ Setup guide: code.claude.com/docs/en/skills

Claude Desktop (App)

Open Settings → Skills → Upload skill folder. Select the image-prompt-engineer-skill folder from the ZIP. The skill appears in your Skills list and stays available across conversations.

→ Setup guide: support.claude.com/en/articles/12512180-using-skills-in-claude

OpenClaw

Copy the skill folder to your OpenClaw skills directory (e.g. `~/.agents/skills/`). Add an entry in your `openclaw.json` config under `skills.entries` with `enabled: true`. Restart the gateway.

→ Setup guide: docs.openclaw.ai/tools/skills

ChatGPT

Two options: (1) Upload SKILL.md as an attachment in any conversation — tell ChatGPT ‘Use the attached skill to optimize my image prompts.’ (2) For persistent use, create a Custom GPT: go to Explore GPTs → Create → paste the contents of SKILL.md into Instructions. This gives you a reusable, shareable prompt engineer.

→ Setup guide: help.openai.com/en/articles/8554397-creating-a-gpt

Google Gemini

Upload SKILL.md as a context document in any Gemini conversation, or create a Gem (Gemini's custom instructions): go to Gem Manager → New Gem → paste SKILL.md contents into the instructions field.

→ Setup guide: support.google.com/gemini/answer/15235603


Any Other LLM

The skill is plain markdown — it works with any LLM that accepts file uploads or system instructions. Upload SKILL.md and the references/ folder, or paste the contents as context. Works with Copilot, Perplexity, local models, and more.

What you get for every image description:

- An **optimized prompt** — with real photography parameters (film stock, lens, lighting), copy-paste ready
- A **model recommendation** — which AI model fits your specific needs
- **Anti-AI realism recipes** — per-model techniques for authentically photographic results
- **Provider options** — where to generate, with estimated pricing
- A **quality score (0–10)** — weighted across 6 dimensions with model-specific nuances
- **What-to-avoid warnings** — model-specific pitfalls that waste generation credits

The optimized prompt is always written in plain English. Paste it into Midjourney, Canva AI, Leonardo, ChatGPT, or any API.



Ready to See AI-Powered Production at Scale?

This toolkit is one piece of a 40+ agent system we built for creative production — video, music, design, and strategy, all working together. The same photography knowledge that took professionals years to learn, mapped to 6 AI models.

We don't compete with your agency — we complete it.

[Schedule a Discovery Call →](#)

inspiredcreativegroupinc.com/partners