
Generate Better AI Images in Minutes

The Prompt Engineering Toolkit

A practical guide for agencies and creative teams.

6 AI Models · Optimized Prompts · Zero Guesswork



inspiredcreativegroupinc.com

WHAT'S INSIDE

This Toolkit

This guide and the **Image Prompt Engineer** skill work together. The skill does the heavy lifting — the guide tells you what to aim at.

■ The Skill	■ This Guide
A Claude skill that transforms plain English into production-ready prompts for 6 AI image models. Handles model selection, keyword optimization, quality scoring, and technical specs automatically.	Covers which model to use, when, and why. Where to generate — from free web UIs to API pipelines. Plus 5 prompt engineering techniques that make an immediate difference.

How they work together:

1. Describe the image you want in plain English.
2. The skill optimizes your prompt for the best model.
3. Take the output to any generation service — web UI or API.
4. This guide helps you pick the right model and service for your workflow.

The skill does **not** generate images — it engineers the best possible prompt. You're never locked into a single provider.

MODEL SELECTION

Why These 6 Models

We evaluated dozens of AI image generators across photorealism, text rendering, design quality, speed, resolution, editing, and price. These 6 cover every professional use case — no gaps, no redundancy.

Model	What It Covers
Flux 2	Photorealism, multi-reference consistency, detailed scenes
Nano Banana Pro	High-resolution (up to 4K), general purpose, natural language
Recraft V4	Vector/SVG output, design assets, marketing materials
Ideogram 3.0	Text rendering, color-controlled branding, style codes
GPT Image 1.5	Instruction following, conversational editing, iteration
Grok Imagine	Quick aesthetics, budget batch work, phone-screen ratios

Each model is best-in-class at something no other does as well. Flux 2 owns photorealism. Recraft V4 is the only model producing real SVG vectors. Ideogram 3.0 leads in text rendering. GPT Image 1.5 is uniquely conversational. Together: full coverage from \$0.02 drafts to production 4K prints.

THE MODELS

6 Models, Explained

MODEL 1 OF 6

Flux 2

by Black Forest Labs

Best for: Photorealistic product shots, portraits, detailed scenes, multi-reference consistency

Key Strengths

- Best photorealism of any current model — natural skin, lighting, materials
- 32K token prompt context — more detail always helps
- Up to 10 reference images for character/style consistency across campaigns
- Sub-10-second generation (Pro/Flex variants)
- Strong text rendering, especially with the Flex variant

Copy-Paste (Web UI)	API (Automation)
Leonardo.ai	Fal.ai — ~\$0.025/image
Freepik AI Image Generator	BFL.ai (direct) — ~\$0.025/image
Any interface offering Flux 2	Replicate — usage-based

Price: \$0.025–\$0.05/image depending on variant and provider

■ *Enable prompt expansion (on by default) — the model enhances your prompt with its own knowledge. For product photography, provide reference images for consistency across an entire catalog.*

MODEL 2 OF 6

Nano Banana Pro

by Google (Gemini 3 Pro Image)

Best for: High-resolution output (up to 4K), general purpose, natural language prompting

Key Strengths

- Up to 4K resolution — ideal for large-format prints and gallery work
- Simple, conversational prompting — describe naturally, skip the keywords
- Google's multimodal intelligence understands context and composition
- Great all-rounder for teams that want one reliable default model

Copy-Paste (Web UI)	API (Automation)
Google AI Studio (Gemini)	Fal.ai — ~\$0.02/image
Any Gemini-powered creative tool	Google AI API — direct Gemini access

Price: ~\$0.02/image

■ Write prompts like you're describing the image to a colleague — conversational language outperforms keyword stacking with this model. Use the '4K' resolution tier for anything going to print.

MODEL 3 OF 6

Recraft V4

by Recraft

Best for: Logos, icons, vector graphics (SVG), design assets, marketing materials

Key Strengths

- The ONLY model producing real, editable SVG vector files
- Outputs: SVG, PNG, JPG, PDF, TIFF, and Lottie animations
- Design-forward aesthetic — understands hierarchy and composition
- Exploration mode: multiple design directions from one prompt
- Free plan available on recraft.ai

Copy-Paste (Web UI)	API (Automation)
Recraft.ai — free web interface	Fal.ai — pricing TBD
Any tool with Recraft integration	Recraft.ai API — free plan available
	Replicate — usage-based

Price: Free plan available; API pricing varies

■ *Use design vocabulary: ‘balanced composition,’ ‘cohesive color palette,’ ‘visual hierarchy.’ Recraft was built for designers and responds to design language better than photography terms.*

MODEL 4 OF 6

Ideogram 3.0

by Ideogram

Best for: Text-in-image rendering, branded content, color-controlled designs, reproducible styles

Key Strengths

- Gold standard for text rendering — signs, logos, titles, dense typography
- Style codes (8-char hex) for exact style reproduction across campaigns
- Color palette with hex weights for brand-precise color control
- Negative prompts to explicitly exclude unwanted elements

- Magic prompt auto-enhancement for quick generation

Copy-Paste (Web UI)	API (Automation)
Ideogram.ai — free web interface	Fal.ai — ~\$0.02/image
ChatGPT (for text-heavy requests)	Ideogram.ai API — direct
Any tool with Ideogram integration	Together.ai
	Replicate

Price: ~\$0.02/image; free tier on ideogram.ai

■ Always put text you want rendered in quotes within the prompt — "SALE 50% OFF". Save your style codes and reuse them across an entire campaign for perfect visual consistency.

MODEL 5 OF 6

GPT Image 1.5

by OpenAI

Best for: Instruction following, multi-turn editing, conversational refinement, concept iteration

Key Strengths

- Best instruction following — understands complex, nuanced directions
- Unique multi-turn editing: generate, then refine ('make the sky more orange')
- Broad real-world knowledge — references brands, locations, styles accurately
- Three quality tiers: low (\$0.009), medium, high (\$0.20) — budget flexibility
- Transparent background support

Copy-Paste (Web UI)	API (Automation)
ChatGPT — native, multi-turn editing	OpenAI API — \$0.009–\$0.20/image
OpenAI Playground	

Price: \$0.009 (low) to \$0.20 (high) per image

■ *Structure your first prompt with clearly separate elements — this makes follow-up edits easier. 'A woman in a red dress, standing in a modern office, with city skyline through windows' lets you later say 'change the dress to blue' without regenerating everything.*

MODEL 6 OF 6

Grok Imagine

by xAI

Best for: Quick aesthetic images, budget batch work, phone wallpapers, exploration

Key Strengths

- Cheapest option at \$0.02/image — great for exploration and drafts
- Unique phone-screen ratios (19.5:9, 20:9) for mobile content
- Fast generation with attractive aesthetic output
- Combined image generation + editing in one API family

Copy-Paste (Web UI)	API (Automation)
Grok (x.ai / X platform)	Fal.ai — \$0.02/image
Any tool with Grok integration	xAI API — direct access

Price: ~\$0.02/image

-
- Focus on mood and aesthetics rather than technical details. Perfect for quick mood boards before committing to a detailed generation with Flux 2 or Ideogram.

AT A GLANCE

Comparison Matrix

	Flux 2	Nano Banana	Recraft V4	Ideogram 3.0	GPT Image	Grok Imagine
Photo-realism	★★★★★	★★★★	★★★★	★★★★	★★★★	★★★
Text	★★★★★	★★★	★★★★	★★★★★	★★★★★	★★★
Speed	★★★★★	★★★★	★★★	★★★★	★★★	★★★★★
Max Res	4MP	4K	2048 ²	Varies	1536px	Varies
SVG	—	—	✓	—	—	—
Editing	★★★★	★	★	★★★★	★★★★★	★★★
Price	\$0.025+	~\$0.02	Free+	~\$0.02	\$0.009+	~\$0.02

Quick Decision Guide

Need a realistic photo? → Flux 2

Need text in the image? → Ideogram 3.0 or GPT Image 1.5

Need a logo or vector? → Recraft V4

Need 4K for print? → Nano Banana Pro

Need to iterate and edit? → GPT Image 1.5

Need it fast and cheap? → Grok Imagine (\$0.02)

Need brand consistency? → Flux 2 (references) or Ideogram (style codes)

5 Prompt Tips That Work

1 Be Specific About What Makes It Good

Don't say 'beautiful sunset.' Say 'golden hour sunset with warm amber light casting long shadows across wet sand, scattered clouds lit from below in pink and coral, gentle wave foam in the foreground.' The more specific, the less the AI guesses — and guessing is where quality drops.

2 Front-Load the Subject

Put your main subject first, then add context. AI models weight the beginning of prompts more heavily. 'A golden retriever sitting in a sunlit meadow, wildflowers in soft focus, photorealistic' works better than burying the subject mid-sentence.

3 Specify Lighting — Always

Lighting is the single biggest quality lever in AI image generation. 'Studio lighting with soft shadows' produces dramatically better results than leaving it unspecified. Good defaults: golden hour, studio lighting, natural window light, dramatic rim lighting.

4 Match Language to the Model

Each model responds to different styles. Flux 2 rewards detail and photography terms. Nano Banana Pro and GPT Image prefer conversation. Recraft V4 responds to design vocabulary. Grok wants aesthetic and mood keywords. The skill handles this automatically.

5 Use the Negative (When Available)

Ideogram 3.0 supports negative prompts — telling the AI what to avoid. 'Negative: blurry, low quality, distorted faces, cluttered background' prevents common failure modes. Even for other models, adding 'avoid blurry details' to the end helps.

How to Use the Skill

The Image Prompt Engineer skill works in any Claude environment. Drop it in and start describing images.

Claude Code (Terminal)

Add the skill folder to your workspace's skills directory. Claude picks it up automatically.

Claude Desktop

Add the skill folder as a project in settings. Include SKILL.md and the references/ folder.

Claude Web ([claude.ai](#))

Upload SKILL.md as an attachment. For deeper model knowledge, also upload files from references/.

What you get for every image description:

- An **optimized prompt** — human-readable, copy-paste ready for any interface
- A **model recommendation** — which AI model fits your specific needs
- **Provider options** — where to generate, with estimated pricing
- A **quality score** (0–10) — how effective the prompt is likely to be
- **Technical specs** — resolution, aspect ratio, and parameters tuned for your use case
- **Warnings** — anything that might cause issues before you generate

The optimized prompt is always written in plain English. Paste it into Midjourney, Canva AI, Leonardo, ChatGPT, or any API.

Ready to See AI-Powered Production at Scale?

This toolkit is one piece of a 40+ agent system we built for creative production — video, music, design, and strategy, all working together.

We don't compete with your agency — we complete it.

[Schedule a Discovery Call →](#)

inspiredcreativegroupinc.com/partners